

Fast Iterative Shrinkage-Thresholding Algorithm (FISTA): A Comprehensive Study of Accelerated Proximal Gradient Methods

Lecture Notes

July 17, 2025

Contents

1	Introduction and Motivation	2
1.1	Overview of Sparse Optimization	2
1.2	The ℓ_1 -Regularized Least Squares Problem	2
1.3	Comparison of Regularization Norms	2
2	Proximal Gradient Methods	3
2.1	Limitations of Classical Gradient Descent	3
2.2	The Proximal Mapping	3
2.3	Proximal Operator for ℓ_1 Regularization	3
2.4	Visualization of Thresholding Operators	4
3	The Proximal Gradient Algorithm	5
3.1	Algorithm Development	5
3.2	ISTA for ℓ_1 -Regularized Least Squares	5
3.3	Step Size Selection	5
3.3.1	Lipschitz Constant Computation	5
3.3.2	Backtracking Line Search	6
3.4	Convergence Analysis	6
4	Fast Iterative Shrinkage-Thresholding Algorithm (FISTA)	7
4.1	Motivation for Acceleration	7
4.2	The FISTA Algorithm	7
4.3	Key Innovations in FISTA	7
4.3.1	The Momentum Sequence	7
4.3.2	The Extrapolation Step	8
4.4	Convergence Theory	8
4.5	Geometric Interpretation	8

5	Implementation Considerations and Extensions	9
5.1	Practical Implementation Details	9
5.1.1	Stopping Criteria	9
5.1.2	Computational Complexity	9
5.2	Extensions and Variants	9
5.2.1	Adaptive Restart	9
5.2.2	Strong Convexity	9
5.3	Applications Beyond ℓ_1 Regularization	10
5.4	Numerical Experiments and Convergence Behavior	10

1 Introduction and Motivation

1.1 Overview of Sparse Optimization

The Fast Iterative Shrinkage-Thresholding Algorithm (FISTA) represents a significant advancement in solving composite optimization problems, particularly those involving sparsity-inducing regularizers. This document provides a comprehensive treatment of the theoretical foundations, algorithmic development, and practical implementation of FISTA.

The fundamental optimization problem we consider takes the form:

$$\min_{\mathbf{x} \in \mathbb{R}^n} F(\mathbf{x}) = f(\mathbf{x}) + g(\mathbf{x}) \quad (1)$$

where:

- $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is a smooth, convex function with Lipschitz continuous gradient
- $g : \mathbb{R}^n \rightarrow \mathbb{R}$ is a convex, possibly non-smooth regularization term
- The composite function F captures both data fidelity and structural constraints

1.2 The ℓ_1 -Regularized Least Squares Problem

A canonical instance of (1) arises in sparse signal recovery and compressed sensing:

$$\min_{\mathbf{x} \in \mathbb{R}^n} \left\{ \frac{1}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2 + \lambda \|\mathbf{x}\|_1 \right\} \quad (2)$$

Here, the objective function comprises:

1. **Data fidelity term:** $f(\mathbf{x}) = \frac{1}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2$
 - $\mathbf{A} \in \mathbb{R}^{m \times n}$ represents the measurement or design matrix
 - $\mathbf{b} \in \mathbb{R}^m$ denotes the observed data vector
 - This quadratic term quantifies the discrepancy between model predictions and observations
2. **Regularization term:** $g(\mathbf{x}) = \lambda \|\mathbf{x}\|_1 = \lambda \sum_{i=1}^n |x_i|$
 - $\lambda > 0$ is the regularization parameter controlling sparsity
 - The ℓ_1 norm promotes sparse solutions by encouraging many components to be exactly zero

1.3 Comparison of Regularization Norms

Norm	Definition	Properties	Optimization	Applications
ℓ_0	$\ \mathbf{x}\ _0 = \{i : x_i \neq 0\} $	Non-convex, discontinuous	NP-hard	Exact sparsity
ℓ_1	$\ \mathbf{x}\ _1 = \sum_i x_i $	Convex, non-smooth	Tractable	Convex relaxation
ℓ_2	$\ \mathbf{x}\ _2 = \sqrt{\sum_i x_i^2}$	Convex, smooth	Closed-form	Ridge regression

Table 1: Comparison of commonly used regularization norms in optimization

2 Proximal Gradient Methods

2.1 Limitations of Classical Gradient Descent

The non-smoothness of the ℓ_1 norm presents a fundamental challenge for classical optimization methods. Consider the subdifferential of the ℓ_1 norm at a point \mathbf{x} :

$$\partial \|\mathbf{x}\|_1 = \left\{ \mathbf{v} \in \mathbb{R}^n : v_i \in \begin{cases} \{1\} & \text{if } x_i > 0 \\ \{-1\} & \text{if } x_i < 0 \\ [-1, 1] & \text{if } x_i = 0 \end{cases} \right\} \quad (3)$$

The multi-valued nature of the subdifferential at $x_i = 0$ precludes the direct application of gradient descent, necessitating more sophisticated approaches.

2.2 The Proximal Mapping

Definition 2.1 (Proximal Operator). *For a convex function $h : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$, the proximal operator is defined as:*

$$\text{prox}_h(\mathbf{v}) = \arg \min_{\mathbf{u} \in \mathbb{R}^n} \left\{ h(\mathbf{u}) + \frac{1}{2} \|\mathbf{u} - \mathbf{v}\|_2^2 \right\} \quad (4)$$

The proximal operator can be interpreted as:

- A generalization of orthogonal projection onto convex sets
- A trade-off between minimizing h and staying close to \mathbf{v}
- An implicit gradient step that handles non-smoothness

2.3 Proximal Operator for ℓ_1 Regularization

Theorem 2.2 (Soft Thresholding). *The proximal operator of $h(\mathbf{x}) = \lambda \|\mathbf{x}\|_1$ is given component-wise by:*

$$[\text{prox}_{\lambda \|\cdot\|_1}(\mathbf{v})]_i = S_\lambda(v_i) = \text{sign}(v_i) \max\{|v_i| - \lambda, 0\} \quad (5)$$

where S_λ is the soft-thresholding operator.

Proof. For the scalar case, we need to solve:

$$\min_{u \in \mathbb{R}} \left\{ \lambda |u| + \frac{1}{2}(u - v)^2 \right\} \quad (6)$$

Taking the subdifferential and setting it to contain zero:

$$0 \in \lambda \cdot \partial |u| + (u - v) \quad (7)$$

$$v \in u + \lambda \cdot \partial |u| \quad (8)$$

Case analysis yields:

- If $v > \lambda$: $u = v - \lambda$
- If $v < -\lambda$: $u = v + \lambda$
- If $|v| \leq \lambda$: $u = 0$

Combining these cases gives the soft-thresholding formula. □

2.4 Visualization of Thresholding Operators

The soft-thresholding operator exhibits the following characteristics:

- **Shrinkage effect:** Non-zero coefficients are reduced by λ
- **Sparsification:** Coefficients with $|v_i| \leq \lambda$ are set to zero
- **Sign preservation:** The sign of large coefficients is maintained

Key Insight: Soft thresholding simultaneously promotes sparsity and shrinks large coefficients

3 The Proximal Gradient Algorithm

3.1 Algorithm Development

The proximal gradient method, also known as the Iterative Shrinkage-Thresholding Algorithm (ISTA), combines gradient descent for the smooth part with proximal operations for the non-smooth part.

Theorem 3.1 (Proximal Gradient Iteration). *For problem (1) with f having L -Lipschitz gradient, the iteration:*

$$\mathbf{x}^{k+1} = \text{prox}_{\alpha g}(\mathbf{x}^k - \alpha \nabla f(\mathbf{x}^k)) \quad (9)$$

converges to the optimal solution when $0 < \alpha < 2/L$.

3.2 ISTA for ℓ_1 -Regularized Least Squares

For the specific problem (2), the algorithm takes the form:

Algorithm 1 Iterative Shrinkage-Thresholding Algorithm (ISTA)

1. **Initialize:** Choose $\mathbf{x}^0 \in \mathbb{R}^n$, step size $\alpha > 0$
 2. **For** $k = 0, 1, 2, \dots$ **do:**
 - (a) Compute gradient: $\mathbf{g}^k = \mathbf{A}^T(\mathbf{A}\mathbf{x}^k - \mathbf{b})$
 - (b) Gradient step: $\mathbf{z}^k = \mathbf{x}^k - \alpha \mathbf{g}^k$
 - (c) Soft threshold: $\mathbf{x}^{k+1} = S_{\alpha\lambda}(\mathbf{z}^k)$
 3. **Until** convergence criterion is met
-

3.3 Step Size Selection

3.3.1 Lipschitz Constant Computation

Definition 3.2 (Lipschitz Continuity of Gradient). *A function f has L -Lipschitz continuous gradient if:*

$$\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\|_2 \leq L \|\mathbf{x} - \mathbf{y}\|_2, \quad \forall \mathbf{x}, \mathbf{y} \in \mathbb{R}^n \quad (10)$$

For the quadratic function $f(\mathbf{x}) = \frac{1}{2} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2^2$:

Lemma 3.3. *The Lipschitz constant of ∇f is $L = \|\mathbf{A}^T \mathbf{A}\|_2 = \lambda_{\max}(\mathbf{A}^T \mathbf{A})$, where λ_{\max} denotes the largest eigenvalue.*

Proof. The gradient is $\nabla f(\mathbf{x}) = \mathbf{A}^T(\mathbf{A}\mathbf{x} - \mathbf{b})$. Thus:

$$\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\|_2 = \|\mathbf{A}^T \mathbf{A}(\mathbf{x} - \mathbf{y})\|_2 \quad (11)$$

$$\leq \|\mathbf{A}^T \mathbf{A}\|_2 \|\mathbf{x} - \mathbf{y}\|_2 \quad (12)$$

$$= \lambda_{\max}(\mathbf{A}^T \mathbf{A}) \|\mathbf{x} - \mathbf{y}\|_2 \quad (13)$$

where we used the fact that the spectral norm equals the largest eigenvalue for symmetric positive semidefinite matrices. \square

3.3.2 Backtracking Line Search

When computing eigenvalues is impractical, adaptive step size selection via backtracking provides a robust alternative:

Algorithm 2 Backtracking Line Search for Proximal Gradient

1. **Parameters:** $\beta \in (0, 1)$ (typically $\beta = 0.5$), $\eta \in (0, 1)$ (typically $\eta = 0.9$)
 2. **Initialize:** $\alpha = \alpha_0$ (initial guess, e.g., $\alpha_0 = 1$)
 3. **Repeat:**
 - (a) Compute: $\mathbf{x}^+ = \text{prox}_{\alpha g}(\mathbf{x}^k - \alpha \nabla f(\mathbf{x}^k))$
 - (b) **While** $F(\mathbf{x}^+) > F(\mathbf{x}^k)$:
 - Set $\alpha \leftarrow \beta \alpha$
 - Recompute \mathbf{x}^+
 4. **Set:** $\mathbf{x}^{k+1} = \mathbf{x}^+$
-

3.4 Convergence Analysis

Theorem 3.4 (ISTA Convergence Rate). *For the proximal gradient method with constant step size $\alpha = 1/L$, we have:*

$$F(\mathbf{x}^k) - F(\mathbf{x}^*) \leq \frac{L \|\mathbf{x}^0 - \mathbf{x}^*\|_2^2}{2k} \quad (14)$$

where \mathbf{x}^* is an optimal solution.

Important:
ISTA
achieves
 $O(1/k)$
conver-
gence

4 Fast Iterative Shrinkage-Thresholding Algorithm (FISTA)

4.1 Motivation for Acceleration

While ISTA achieves $O(1/k)$ convergence, Nesterov's acceleration technique can improve this to $O(1/k^2)$ without additional computational cost per iteration. This acceleration is achieved through a momentum-like mechanism that exploits the history of iterates.

4.2 The FISTA Algorithm

Algorithm 3 Fast Iterative Shrinkage-Thresholding Algorithm (FISTA)

1. **Initialize:**

- Choose $\mathbf{x}^0 = \mathbf{y}^1 \in \mathbb{R}^n$
- Set $t_1 = 1$
- Choose step size $\alpha \leq 1/L$

2. **For** $k = 1, 2, 3, \dots$ **do:**

(a) Proximal gradient step:

$$\mathbf{x}^k = \text{prox}_{\alpha g}(\mathbf{y}^k - \alpha \nabla f(\mathbf{y}^k)) \quad (15)$$

(b) Update momentum parameter:

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2} \quad (16)$$

(c) Compute extrapolated point:

$$\mathbf{y}^{k+1} = \mathbf{x}^k + \frac{t_k - 1}{t_{k+1}}(\mathbf{x}^k - \mathbf{x}^{k-1}) \quad (17)$$

3. **Until** convergence

4.3 Key Innovations in FISTA

4.3.1 The Momentum Sequence

The sequence $\{t_k\}$ satisfies the recurrence relation:

$$t_{k+1}^2 - t_{k+1} - t_k^2 = 0 \quad (18)$$

This yields the closed-form expression:

$$t_k = \frac{k+1}{2} + O(1) \approx \frac{k}{2} \text{ for large } k \quad (19)$$

4.3.2 The Extrapolation Step

The extrapolation coefficient:

$$\beta_k = \frac{t_k - 1}{t_{k+1}} \approx \frac{k - 2}{k + 1} \rightarrow 1 \text{ as } k \rightarrow \infty \quad (20)$$

This creates an "overshoot" effect that accelerates convergence by anticipating the trajectory of the iterates.

4.4 Convergence Theory

Theorem 4.1 (FISTA Convergence Rate). *For FISTA with step size $\alpha = 1/L$, the following bound holds:*

$$F(\mathbf{x}^k) - F(\mathbf{x}^*) \leq \frac{2L \|\mathbf{x}^0 - \mathbf{x}^*\|_2^2}{(k + 1)^2} \quad (21)$$

Remark 4.2. *The $O(1/k^2)$ rate is optimal for first-order methods on the class of convex functions with Lipschitz continuous gradients.*

4.5 Geometric Interpretation

FISTA can be viewed as performing gradient descent on an auxiliary sequence $\{\mathbf{y}^k\}$ that is constructed to have favorable properties:

- The sequence $\{\mathbf{y}^k\}$ exhibits less oscillation than $\{\mathbf{x}^k\}$
- The extrapolation step creates a "look-ahead" effect
- The momentum builds up over iterations, accelerating convergence in consistent directions

5 Implementation Considerations and Extensions

5.1 Practical Implementation Details

5.1.1 Stopping Criteria

Common convergence criteria for FISTA include:

1. **Relative change in objective:**

$$\frac{|F(\mathbf{x}^k) - F(\mathbf{x}^{k-1})|}{|F(\mathbf{x}^{k-1})|} < \epsilon_{\text{obj}} \quad (22)$$

2. **Relative change in iterates:**

$$\frac{\|\mathbf{x}^k - \mathbf{x}^{k-1}\|_2}{\|\mathbf{x}^{k-1}\|_2} < \epsilon_{\text{sol}} \quad (23)$$

3. **Optimality conditions:**

$$\text{dist}(0, \partial F(\mathbf{x}^k)) < \epsilon_{\text{opt}} \quad (24)$$

5.1.2 Computational Complexity

Per iteration, FISTA requires:

- One gradient evaluation: $O(mn)$ for matrix-vector products
- One soft-thresholding operation: $O(n)$
- Vector operations: $O(n)$

Total complexity: $O(mn)$ per iteration, same as ISTA but with faster convergence.

5.2 Extensions and Variants

5.2.1 Adaptive Restart

Adaptive restart strategies can further improve practical performance:

$$\text{Restart if: } \langle \mathbf{y}^k - \mathbf{x}^k, \mathbf{x}^k - \mathbf{x}^{k-1} \rangle > 0 \quad (25)$$

This condition detects when the momentum is counterproductive.

5.2.2 Strong Convexity

When f is μ -strongly convex, linear convergence can be achieved:

$$F(\mathbf{x}^k) - F(\mathbf{x}^*) \leq \left(1 - \sqrt{\frac{\mu}{L}}\right)^k [F(\mathbf{x}^0) - F(\mathbf{x}^*)] \quad (26)$$

Regularizer	Proximal Operator	Application
$\ \mathbf{x}\ _1$	Soft thresholding	Sparse recovery
$\ \mathbf{x}\ _2$	Scaling	Group sparsity
$\delta_C(\mathbf{x})$	Projection onto C	Constrained optimization
$\ \mathbf{X}\ _*$	Singular value thresholding	Low-rank matrix recovery

Table 2: Common regularizers and their proximal operators

5.3 Applications Beyond ℓ_1 Regularization

FISTA’s framework extends to various proximal operators:

5.4 Numerical Experiments and Convergence Behavior

In practice, FISTA exhibits several characteristic behaviors:

1. **Initial phase:** Rapid decrease in objective value
2. **Middle phase:** Steady convergence with momentum benefits
3. **Final phase:** Oscillations may occur near the solution

Comparison with ISTA Empirical studies consistently show FISTA requiring 5-10% fewer iterations than ISTA for the same accuracy, validating the theoretical acceleration.

Conclusion

The Fast Iterative Shrinkage-Thresholding Algorithm represents a fundamental advancement in composite convex optimization, combining:

- Elegant handling of non-smooth regularizers via proximal operators
- Optimal convergence rates through Nesterov’s acceleration
- Practical efficiency and broad applicability

FISTA’s success has inspired numerous extensions and remains a cornerstone algorithm in machine learning, signal processing, and computational statistics.